

EEL5733/EEL4732 Advanced Systems  
Programming  
Final Exam

Student Name:

UFID:

Signature:

**Important Note** There are four questions. Please read all the questions carefully and write your name on all answer pages. Good luck.

## Questions

1. (7 pts) In a multi-process application, a developer uses the `mmap` system call to let two processes in the application to communicate via shared memory without a backing file. For some reason, the communication does not work, i.e., one process does not see the values written by the other process and vice versa. Clearly describe with code snippets two different scenarios such that in each scenario a different type of programming mistake prevents correctly setting up memory sharing based inter-process communication. Please note that compilation errors would not be accepted. Also, make sure to include the `mmap` call(s) and any other relevant in your code snippets, and briefly explain the mistake for each case. *Hint:* Think about how to properly use the `mmap` system call to enable inter-process communication when there is no backing file and figure out how to introduce the relevant bugs.

(a) Scenario 1

(b) Scenario 2

2. ( 9 pts) This question is about the `short` driver that was discussed in the Interrupt Handling chapter of the Linux Device Drivers book (<https://static.lwn.net/images/pdf/LDD3/ch10.pdf>). The source code of the driver can also be found at <https://github.com/martinezjavier/ldd3/blob/master/short/short.c>.

- (a) What happens if the call to the `ioremap` function at line 21 in the `short_init` function gets commented out? Explain.
- (b) List all line no's that request some kernel resources in the code snippet of the `short_init` function below and briefly explain each kernel resource that you mention.
- (c) List the names of all top handlers as registered by the `short_init` function. Briefly explain how each of the top handler differs from the others.

```
1  int short_init(void)
2  {
3      int result;
4
5      short_base = base;
6      short_irq = irq;
7
8      if (!use_mem) {
9          if (! request_region(short_base, SHORT_NR_PORTS, "short")) {
10             printk(KERN_INFO "short: can't get I/O port address 0x%lx\n",
11                          short_base);
12             return -ENODEV;
13         }
14     } else {
15         if (! request_mem_region(short_base, SHORT_NR_PORTS, "short")) {
16             printk(KERN_INFO "short: can't get I/O mem address 0x%lx\n",
17                          short_base);
18             return -ENODEV;
19         }
20     }
21     short_base = (unsigned long) ioremap(short_base, SHORT_NR_PORTS);
22 }
23 result = register_chrdev(major, "short", &short_fops);
24 if (result < 0) {
25     printk(KERN_INFO "short: can't get major number\n");
26     release_region(short_base, SHORT_NR_PORTS);
27     return result;
28 }
29 if (major == 0) major = result; /* dynamic */
30 short_buffer = __get_free_pages(GFP_KERNEL, 0);
31 short_head = short_tail = short_buffer;
```

```

32
33 INIT_WORK(&short_wq, (void (*)(struct work_struct *)) short_do_tasklet);
34
35
36 if (short_irq < 0 && probe == 1)
37     short_kernelprobe();
38
39 if (short_irq < 0 && probe == 2)
40     short_selfprobe();
41
42 if (short_irq < 0)
43     switch(short_base) {
44         case 0x378: short_irq = 7; break;
45         case 0x278: short_irq = 2; break;
46         case 0x3bc: short_irq = 5; break;
47     }
48
49 if (short_irq >= 0 && share > 0) {
50     result = request_irq(short_irq, short_sh_interrupt,
51                          IRQF_SHARED, "short",
52                          short_sh_interrupt);
53     if (result) {
54         printk(KERN_INFO "short: can't get assigned irq %i\n", short_irq);
55         short_irq = -1;
56     }
57     else {
58         outb(0x10, short_base+2);
59     }
60     return 0;
61 }
62
63 if (short_irq >= 0) {
64     result = request_irq(short_irq, short_interrupt,
65                          0, "short", NULL);
66     if (result) {
67         printk(KERN_INFO "short: can't get assigned irq %i\n",
68                  short_irq);
69         short_irq = -1;
70     }
71     else {
72         outb(0x10, short_base+2);
73     }
74 }
75
76 if (short_irq >= 0 && (wq + tasklet) > 0) {
77     free_irq(short_irq, NULL);
78     result = request_irq(short_irq,
79                          tasklet ? short_tl_interrupt :
80                          short_wq_interrupt,

```

```
81                                     0, "short-bh", NULL);
82     if (result) {
83         printk(KERN_INFO "short-bh: can't get assigned irq %i\n",
84                        short_irq);
85         short_irq = -1;
86     }
87 }
88 return 0;
89 }
```

3. (9 pts) This question is about the `scull` driver. You can find the full source code at

<https://github.com/martinezjavier/ldd3/tree/master/scull>.

- (a) In the `scull_init_module` function, what happens if `scull_setup_cdev(&scull_devices[i], i);` is executed as the first statement of the body of the `for` loop, e.g, switch the statements at lines 38 and 41?
- (b) What is the role of the `container_of` function that is used in the `scull_open` function (line 67)? Explain.
- (c) What would be the side-effect of commenting out the assignment statement `filp->private_data = dev;` at line 68 in the `scull_open` function? Explain.

```
1  struct scull_dev {
2      struct scull_qset *data; /* Pointer to first quantum set */
3      int quantum;             /* the current quantum size */
4      int qset;                /* the current array size */
5      unsigned long size;      /* amount of data stored here */
6      unsigned int access_key; /* used by sculluid and scullpriv */
7      struct mutex lock;       /* mutual exclusion semaphore */
8      struct cdev cdev;        /* Char device structure */
9  };
10
11 int scull_init_module(void)
12 {
13     int result, i;
14     dev_t dev = 0;
15
16     if (scull_major) {
17         dev = MKDEV(scull_major, scull_minor);
18         result = register_chrdev_region(dev, scull_nr_devs, "scull");
19     } else {
20         result = alloc_chrdev_region(&dev, scull_minor, scull_nr_devs,
21                                     "scull");
22         scull_major = MAJOR(dev);
23     }
24     if (result < 0) {
25         printk(KERN_WARNING "scull: can't get major %d\n", scull_major);
26         return result;
27     }
28
29     scull_devices = kmalloc(scull_nr_devs * sizeof(struct scull_dev), GFP_KERNEL);
30     if (!scull_devices) {
31         result = -ENOMEM;
32         goto fail; /* Make this more graceful */

```

```

33     }
34     memset(scull_devices, 0, scull_nr_devs * sizeof(struct scull_dev));
35
36     /* Initialize each device. */
37     for (i = 0; i < scull_nr_devs; i++) {
38         scull_devices[i].quantum = scull_quantum;
39         scull_devices[i].qset = scull_qset;
40         mutex_init(&scull_devices[i].lock);
41         scull_setup_cdev(&scull_devices[i], i);
42     }
43
44     return 0; /* succeed */
45
46 fail:
47     scull_cleanup_module();
48     return result;
49 }
50
51 static void scull_setup_cdev(struct scull_dev *dev, int index)
52 {
53     int err, devno = MKDEV(scull_major, scull_minor + index);
54
55     cdev_init(&dev->cdev, &scull_fops);
56     dev->cdev.owner = THIS_MODULE;
57     dev->cdev.ops = &scull_fops;
58     err = cdev_add (&dev->cdev, devno, 1);
59     if (err)
60         printk(KERN_NOTICE "Error %d adding scull%d", err, index);
61 }
62
63 int scull_open(struct inode *inode, struct file *filp)
64 {
65     struct scull_dev *dev;
66
67     dev = container_of(inode->i_cdev, struct scull_dev, cdev);
68     filp->private_data = dev;
69
70     if ( (filp->f_flags & O_ACCMODE) == O_WRONLY) {
71         if (mutex_lock_interruptible(&dev->lock))
72             return -ERESTARTSYS;
73         scull_trim(dev); /* ignore errors */
74         mutex_unlock(&dev->lock);
75     }
76     return 0;
77 }
78
79 loff_t scull_llseek(struct file *filp, loff_t off, int whence)
80 {
81     struct scull_dev *dev = filp->private_data;

```

```

82     loff_t newpos;
83
84     switch(whence) {
85         case 0: /* SEEK_SET */
86             newpos = off;
87             break;
88
89         case 1: /* SEEK_CUR */
90             newpos = filp->f_pos + off;
91             break;
92
93         case 2: /* SEEK_END */
94             newpos = dev->size + off;
95             break;
96
97         default: /* can't happen */
98             return -EINVAL;
99     }
100     if (newpos < 0) return -EINVAL;
101     filp->f_pos = newpos;
102     return newpos;
103 }

```



4. (10 pts) This question is about the `usbkbd` driver that we studied in class. The source code is provided with the line numbers below for your reference.
- (a) What would be the consequence of commenting out the call to the `usb_submit_urb` function in the `usb_kbd_open` function (lines 236 and 237)? Explain in terms of the `usb_kbd_irq` and the `usb_submit_led` functions.
  - (b) Which buffer(s) in the driver may get accessed by the keyboard by transferring data to or transferring data from? Explain how you identify those buffer(s) and the type of operation (read or write) that gets performed by the device for each such buffer that you identify.
  - (c) Why does the `usbkbd` driver does not define any Virtual File System entry points? Does it mean the keyboard device cannot be accessed through the `open` system call from the user space? Explain.
  - (d) Specify the line number from which the very first LED urb get submitted. Also, explain what gets achieved with the submission of an LED urb?

```

93 struct usb_kbd {
94     struct input_dev *dev;
95     struct usb_device *usbdev;
96     unsigned char old[8];
97     struct urb *irq, *led;
98     unsigned char newleds;
99     char name[128];
100    char phys[64];
101
102    unsigned char *new;
103    struct usb_ctrlrequest *cr;
104    unsigned char *leds;
105    dma_addr_t new_dma;
106    dma_addr_t leds_dma;
107
108    spinlock_t leds_lock;
109    bool led_urb_submitted;
110
111 };
112
113 static void usb_kbd_irq(struct urb *urb)
114 {
115     struct usb_kbd *kbd = urb->context;
116     int i;
117
118     switch (urb->status) {
119     case 0:                /* success */
120         break;
121     case -ECONNRESET:      /* unlink */
122     case -ENOENT:
123     case -ESHUTDOWN:
124         return;
125     /* -EPIPE: should clear the halt */
126     default:                /* error */
127         goto resubmit;
128     }
129
130     for (i = 0; i < 8; i++)
131         input_report_key(kbd->dev, usb_kbd_keycode[i + 224],
132                         (kbd->new[0] >> i) & 1);
133
134     for (i = 2; i < 8; i++) {
135         if (kbd->old[i] > 3 &&
136             memscan(kbd->new + 2, kbd->old[i], 6) == kbd->new + 8) {
137             if (usb_kbd_keycode[kbd->old[i]])
138                 input_report_key(kbd->dev,
139                                 usb_kbd_keycode[kbd->old[i]], 0);
140             else

```

```

139             hid_info(urb->dev,
140                       "Unknown key (scancode %#x) released.\n",
141                       kbd->old[i]);
142         }
143
144         if (kbd->new[i] > 3 && memscan(kbd->old + 2,
145                                     kbd->new[i], 6) == kbd->old + 8) {
146             if (usb_kbd_keycode[kbd->new[i]])
147                 input_report_key(kbd->dev,
148                                 usb_kbd_keycode[kbd->new[i]], 1);
149             else
150                 hid_info(urb->dev,
151                         "Unknown key (scancode %#x) pressed.\n",
152                         kbd->new[i]);
153         }
154     }
155
156     input_sync(kbd->dev);
157
158     memcpy(kbd->old, kbd->new, 8);
159
160 resubmit:
161     i = usb_submit_urb (urb, GFP_ATOMIC);
162     if (i)
163         hid_err(urb->dev, "can't resubmit intr, %s-%s/input0, status %d",
164                kbd->usbdev->bus->bus_name,
165                kbd->usbdev->devpath, i);
166 }
167
168 static int usb_kbd_event(struct input_dev *dev, unsigned int type,
169                          unsigned int code, int value)
170 {
171     unsigned long flags;
172     struct usb_kbd *kbd = input_get_drvdata(dev);
173
174     if (type != EV_LED)
175         return -1;
176
177     spin_lock_irqsave(&kbd->leds_lock, flags);
178     kbd->newleds = (!test_bit(LED_KANA, dev->led) << 3) |
179                  (!test_bit(LED_COMPOSE, dev->led) << 3) |
180                  (!test_bit(LED_SCROLLL, dev->led) << 2) |
181                  (!test_bit(LED_CAPSL, dev->led) << 1) |
182                  (!test_bit(LED_NUML, dev->led));
183
184     if (kbd->led_urb_submitted){
185         spin_unlock_irqrestore(&kbd->leds_lock, flags);
186         return 0;
187     }

```

```

184
185     if (*(kbd->leds) == kbd->newleds){
186         spin_unlock_irqrestore(&kbd->leds_lock, flags);
187         return 0;
188     }
189
190     *(kbd->leds) = kbd->newleds;
191
192     kbd->led->dev = kbd->usbdev;
193     if (usb_submit_urb(kbd->led, GFP_ATOMIC))
194         pr_err("usb_submit_urb(leds) failed\n");
195     else
196         kbd->led_urb_submitted = true;
197
198     spin_unlock_irqrestore(&kbd->leds_lock, flags);
199
200     return 0;
201 }
202
203 static void usb_kbd_led(struct urb *urb)
204 {
205     unsigned long flags;
206     struct usb_kbd *kbd = urb->context;
207
208     if (urb->status)
209         hid_warn(urb->dev, "led urb status %d received\n",
210                 urb->status);
211
212     spin_lock_irqsave(&kbd->leds_lock, flags);
213
214     if (*(kbd->leds) == kbd->newleds){
215         kbd->led_urb_submitted = false;
216         spin_unlock_irqrestore(&kbd->leds_lock, flags);
217         return;
218     }
219
220     *(kbd->leds) = kbd->newleds;
221
222     kbd->led->dev = kbd->usbdev;
223     if (usb_submit_urb(kbd->led, GFP_ATOMIC)){
224         hid_err(urb->dev, "usb_submit_urb(leds) failed\n");
225         kbd->led_urb_submitted = false;
226     }
227     spin_unlock_irqrestore(&kbd->leds_lock, flags);
228
229 }
230
231 static int usb_kbd_open(struct input_dev *dev)
232 {

```

```

233     struct usb_kbd *kbd = input_get_drvdata(dev);
234
235     kbd->irq->dev = kbd->usbdev;
236     if (usb_submit_urb(kbd->irq, GFP_KERNEL))
237         return -EIO;
238
239     return 0;
240 }
241
242 static void usb_kbd_close(struct input_dev *dev)
243 {
244     struct usb_kbd *kbd = input_get_drvdata(dev);
245
246     usb_kill_urb(kbd->irq);
247 }
248
249 static int usb_kbd_alloc_mem(struct usb_device *dev, struct usb_kbd *kbd)
250 {
251     if (!(kbd->irq = usb_alloc_urb(0, GFP_KERNEL)))
252         return -1;
253     if (!(kbd->led = usb_alloc_urb(0, GFP_KERNEL)))
254         return -1;
255     if (!(kbd->new = usb_alloc_coherent(dev, 8, GFP_ATOMIC, &kbd->new_dma)))
256         return -1;
257     if (!(kbd->cr = kmalloc(sizeof(struct usb_ctrlrequest), GFP_KERNEL)))
258         return -1;
259     if (!(kbd->leds = usb_alloc_coherent(dev, 1, GFP_ATOMIC, &kbd->leds_dma)))
260         return -1;
261
262     return 0;
263 }
264
265 static void usb_kbd_free_mem(struct usb_device *dev, struct usb_kbd *kbd)
266 {
267     usb_free_urb(kbd->irq);
268     usb_free_urb(kbd->led);
269     usb_free_coherent(dev, 8, kbd->new, kbd->new_dma);
270     kfree(kbd->cr);
271     usb_free_coherent(dev, 1, kbd->leds, kbd->leds_dma);
272 }
273
274 static int usb_kbd_probe(struct usb_interface *iface,
275                          const struct usb_device_id *id)
276 {
277     struct usb_device *dev = interface_to_usbdev(iface);
278     struct usb_host_interface *interface;
279     struct usb_endpoint_descriptor *endpoint;
280     struct usb_kbd *kbd;
281     struct input_dev *input_dev;

```

```

282     int i, pipe, maxp;
283     int error = -ENOMEM;
284
285     interface = iface->cur_altsetting;
286
287     if (interface->desc.bNumEndpoints != 1)
288         return -ENODEV;
289
290     endpoint = &interface->endpoint[0].desc;
291     if (!usb_endpoint_is_int_in(endpoint))
292         return -ENODEV;
293
294     pipe = usb_rcvintpipe(dev, endpoint->bEndpointAddress);
295     maxp = usb_maxpacket(dev, pipe, usb_pipeout(pipe));
296
297     kbd = kzalloc(sizeof(struct usb_kbd), GFP_KERNEL);
298     input_dev = input_allocate_device();
299     if (!kbd || !input_dev)
300         goto fail1;
301
302     if (usb_kbd_alloc_mem(dev, kbd))
303         goto fail2;
304
305     kbd->usbdev = dev;
306     kbd->dev = input_dev;
307     spin_lock_init(&kbd->leds_lock);
308
309     if (dev->manufacturer)
310         strlcpy(kbd->name, dev->manufacturer, sizeof(kbd->name));
311
312     if (dev->product) {
313         if (dev->manufacturer)
314             strlcat(kbd->name, " ", sizeof(kbd->name));
315         strlcat(kbd->name, dev->product, sizeof(kbd->name));
316     }
317
318     if (!strlen(kbd->name))
319         snprintf(kbd->name, sizeof(kbd->name),
320                 "USB HIDBP Keyboard %04x:%04x",
321                 le16_to_cpu(dev->descriptor.idVendor),
322                 le16_to_cpu(dev->descriptor.idProduct));
323
324     usb_make_path(dev, kbd->phys, sizeof(kbd->phys));
325     strlcat(kbd->phys, "/input0", sizeof(kbd->phys));
326
327     input_dev->name = kbd->name;
328     input_dev->phys = kbd->phys;
329     usb_to_input_id(dev, &input_dev->id);
330     input_dev->dev.parent = &iface->dev;

```

```

331
332     input_set_drvdata(input_dev, kbd);
333
334     input_dev->evbit[0] = BIT_MASK(EV_KEY) | BIT_MASK(EV_LED) |
335         BIT_MASK(EV_REP);
336     input_dev->ledbit[0] = BIT_MASK(LED_NUML) | BIT_MASK(LED_CAPSL) |
337         BIT_MASK(LED_SCROLLL) | BIT_MASK(LED_COMPOSE) |
338         BIT_MASK(LED_KANA);
339
340     for (i = 0; i < 255; i++)
341         set_bit(usb_kbd_keycode[i], input_dev->keybit);
342     clear_bit(0, input_dev->keybit);
343
344     input_dev->event = usb_kbd_event;
345     input_dev->open = usb_kbd_open;
346     input_dev->close = usb_kbd_close;
347
348     usb_fill_int_urb(kbd->irq, dev, pipe,
349         kbd->new, (maxp > 8 ? 8 : maxp),
350         usb_kbd_irq, kbd, endpoint->bInterval);
351     kbd->irq->transfer_dma = kbd->new_dma;
352     kbd->irq->transfer_flags |= URB_NO_TRANSFER_DMA_MAP;
353
354     kbd->cr->bRequestType = USB_TYPE_CLASS | USB_RECIP_INTERFACE;
355     kbd->cr->bRequest = 0x09;
356     kbd->cr->wValue = cpu_to_le16(0x200);
357     kbd->cr->wIndex = cpu_to_le16(interface->desc.bInterfaceNumber);
358     kbd->cr->wLength = cpu_to_le16(1);
359
360     usb_fill_control_urb(kbd->led, dev, usb_sndctrlpipe(dev, 0),
361         (void *) kbd->cr, kbd->leds, 1,
362         usb_kbd_led, kbd);
363     kbd->led->transfer_dma = kbd->leds_dma;
364     kbd->led->transfer_flags |= URB_NO_TRANSFER_DMA_MAP;
365
366     error = input_register_device(kbd->dev);
367     if (error)
368         goto fail2;
369
370     usb_set_intfdata(iface, kbd);
371     device_set_wakeup_enable(&dev->dev, 1);
372     return 0;
373
374 fail2:
375     usb_kbd_free_mem(dev, kbd);
376 fail1:
377     input_free_device(input_dev);
378     kfree(kbd);
379     return error;

```

```

380 }
381
382 static void usb_kbd_disconnect(struct usb_interface *intf)
383 {
384     struct usb_kbd *kbd = usb_get_intfdata (intf);
385
386     usb_set_intfdata(intf, NULL);
387     if (kbd) {
388         usb_kill_urb(kbd->irq);
389         input_unregister_device(kbd->dev);
390         usb_kill_urb(kbd->led);
391         usb_kbd_free_mem(interface_to_usbdev(intf), kbd);
392         kfree(kbd);
393     }
394 }
395
396 static struct usb_device_id usb_kbd_id_table [] = {
397     { USB_INTERFACE_INFO(USB_INTERFACE_CLASS_HID, USB_INTERFACE_SUBCLASS_BOOT,
398         USB_INTERFACE_PROTOCOL_KEYBOARD) },
399     { } /* Terminating entry */
400 };
401
402 MODULE_DEVICE_TABLE (usb, usb_kbd_id_table);
403
404 static struct usb_driver usb_kbd_driver = {
405     .name = "usbkbd",
406     .probe = usb_kbd_probe,
407     .disconnect = usb_kbd_disconnect,
408     .id_table = usb_kbd_id_table,
409 };
410
411 module_usb_driver(usb_kbd_driver);
412

```